From the Hive to the Hard Drive: A Governed Memory Architecture for Multi-Agent AI Inspired by Biological Swarms

DATE: 2026-06-21

Executive Summary

The development of sophisticated multi-agent artificial intelligence (AI) systems is increasingly constrained by the challenge of creating robust, scalable, and auditable memory architectures. Current systems often struggle with context pollution, semantic drift, and the inability to govern agent actions, leading to unpredictable and untrustworthy behavior [55, 56]. This report presents a focused-depth analysis of collective memory mechanisms in biological swarms to inform the design of a novel, governed multi-agent memory system. We argue that principles evolved over millions of years by social insects, microbial colonies, and migratory animals offer a potent blueprint for engineering resilient and intelligent artificial collectives [31, 76].

The report begins by surveying a diverse range of biological systems—including ants, honeybees, termites, bacterial biofilms, slime molds, the immune system, fungal networks, and migratory animal groups. We analyze how these systems discover, communicate, preserve, update, and discard operational knowledge across individuals, colonies, and generations. Key mechanisms identified include the stigmergic use of the environment as an external memory substrate (ants, termites) [1, 9], the epigenetic reprogramming of innate cells for long-term functional memory (immune system) [29, 30], and the cultural transmission of complex route knowledge in social learners (migratory mammals) [39, 40].

Drawing on these biological insights, we map these natural strategies to modern AI constructs such as Ant Colony Optimization (ACO), blackboard systems, multi-agent reinforcement learning (MARL), and evolutionary algorithms [7, 42, 46]. This analysis reveals a clear disconnect: while biological systems have evolved sophisticated governance and memory decay mechanisms to ensure collective stability, many artificial systems lack the architectural primitives for provenance, auditability, and controlled forgetting [55, 56].

The core contribution of this report is a concrete proposal for a **Governed Multi-Agent Memory Architecture**. This architecture is founded on the principle of separating an agent's free internal cognition from its externally governed actions [53]. Its key components include:
1. **A Structured Memory Schema:** Each memory record is enriched with extensive metadata, including cryptographic provenance, derivation lineage, relevance scores, and verification signatures, making the memory state fully inspectable [48, 55].
2. **A Multi-Signal Retrieval and Scoring Model:** A relevance model that balances semantic similarity, temporal decay, historical utility, and provenance trust to prevent stale or untrustworthy information from polluting an agent's context [52, 55].
3. **A Governed Lifecycle with Provenance:** A robust memory lifecycle that tracks information from ingestion to archival, with consequential external actions requiring a "justification proof" and generating a non-repudiable "receipt," creating a tamper-evident audit trail [48, 50, 53].
4. **A Cross-Generational Transfer Model:** A mechanism for memory inheritance and distillation, inspired by cultural transmission and epigenetic programming, allowing new agent generations to be initialized with curated, high-value knowledge from their predecessors [29, 40, 47].

By translating the time-tested principles of biological swarms into concrete architectural constraints, this report provides a roadmap for building multi-agent AI systems that are not only more intelligent and adaptive but also more transparent, accountable, and safe for deployment in high-stakes, real-world applications.

1. Introduction

The pursuit of artificial general intelligence has increasingly shifted focus from monolithic models to decentralized, collaborative multi-agent systems (MAS) [77]. Inspired by the emergent intelligence of biological swarms, these systems promise unprecedented scalability, robustness, and problem-solving capabilities by distributing tasks among specialized, autonomous agents [76, 78]. However, this decentralized paradigm introduces a fundamental challenge: the management of collective memory. Without a coherent, persistent, and trustworthy memory, a swarm of agents risks devolving into a cacophony of isolated actors, incapable of long-term learning, coordinated action, or stable adaptation [58].

Current memory architectures for LLM-based agents, often extensions of Retrieval-Augmented Generation (RAG), are proving inadequate [59]. They are vulnerable to a host of systemic failures, including context pollution from irrelevant data, semantic drift from repeated summarization, and memory poisoning from malicious inputs [55, 56]. More critically, they lack the governance mechanisms to audit agent reasoning or control external actions, creating a “responsibility gap” that hinders their deployment in mission-critical domains [49, 50]. An agent that cannot explain why it took an action based on its memory is an agent that cannot be trusted.

This report addresses this critical gap by turning to the most successful multi-agent systems known: biological collectives. For eons, social insects, microbial colonies, and even our own immune systems have solved the problem of distributed cognition [31]. They have evolved sophisticated mechanisms to create, share, and manage collective memory, enabling them to navigate complex environments, allocate resources efficiently, and adapt across generations [60]. These biological systems offer a rich, time-tested blueprint for designing the next generation of artificial memory architectures.

The primary objective of this report is to survey these biological memory mechanisms and translate them into a concrete, actionable architectural proposal for a **governed multi-agent memory system**. Our goal is to move beyond simple information storage and retrieval toward a system that embodies principles of provenance, auditability, and principled forgetting [48, 55]. We will explore how ants use the environment as an auditable public ledger [1, 9], how the immune system uses epigenetic reprogramming for long-term adaptation [29], and how migratory herds use cultural transmission to pass vital knowledge to new generations [39, 40]. These insights will form the foundation of an architecture designed to support a system where agents possess freedom of internal cognition but are subject to rigorous external governance for their actions and memory modifications, ensuring the collective remains both intelligent and accountable [53].

2. Methodology

The research for this report was conducted between May and June 2026, with the objective of synthesizing findings from biology, computer science, and AI research to inform the design of a multi-agent memory architecture. The methodology involved a multi-stage process of data collection, analysis, and synthesis.

The primary sources consulted were peer-reviewed academic journals, conference proceedings, and high-quality pre-print archives (e.g., arXiv.org), spanning the fields of entomology, microbiology, immunology, ecology, collective animal behavior, multi-agent systems, artificial intelligence, and cybersecurity. Keyword-driven searches were performed across databases like Google Scholar, PubMed, ScienceDirect, and the ACM Digital Library. Queries included terms such as “collective memory,” “swarm intelligence,” “stigmergy,” “multi-agent memory architecture,” “provenance,” “audit trail,” and “trained immunity” [9, 31, 37, 55, 49].

The analysis phase focused on identifying and extracting core mechanisms of information storage, transmission, decay, and governance within each biological system. A comparative framework was used to map these biological principles to existing or theoretical AI concepts, such as Ant Colony Optimization, blackboard systems, and multi-agent reinforcement learning [7, 42, 46]. This mapping was crucial for bridging the conceptual gap between natural and artificial systems.

The final synthesis phase focused on constructing the proposed “Governed Multi-Agent Memory Architecture.” This involved integrating cutting-edge concepts from recent AI safety and architecture research—such as verifiable memory governance [56], sovereign execution brokers [53], and evidence chains [50]—with the foundational principles derived from the biological survey.

A key limitation of this study is the inherent abstraction required when translating complex biological processes into computational models. Biological realities are often far more nuanced and complex than their algorithmic inspirations. Therefore, this report clearly distinguishes between established biological facts and the speculative architectural recommendations inspired by them, using hedging language where appropriate. The aim is not to replicate biology verbatim but to derive robust, first-principle design patterns for engineering artificial systems.

---

3. Biological Foundations of Collective Memory

Nature is the original architect of decentralized systems. Across vast scales of life, from microscopic bacteria to continent-spanning animal migrations, organisms have evolved sophisticated strategies to pool information and generate collective intelligence [31]. This memory is not always stored in a brain; it can be embedded in the environment, encoded in chemical gradients, programmed into cellular epigenetics, or passed down through culture [9, 29, 39]. Understanding these diverse substrates of collective memory provides a rich foundation for designing artificial systems.

3.1. Social Insects (Ants, Termites, Honeybees, Social wasps)

Social insects represent a paramount example of collective cognition, where colonies of individuals with limited cognitive capacity achieve remarkable feats of engineering, logistics, and defense [79]. Their memory systems are a hybrid of internal (individual) and external (collective) mechanisms.

Ants provide a classic model of stigmergy, a form of indirect communication where individuals interact by modifying their environment [1, 9]. Worker ants, upon discovering a food source, lay down volatile chemical trails of **pheromones**. The trail itself acts as an externalized, collective memory [1]. Its strength, reinforced by returning foragers, encodes the quality and popularity of a food source. The natural **decay** or evaporation of these pheromones is a crucial feature, not a flaw; it functions as a “forgetting” mechanism, ensuring the colony does not remain fixated on depleted resources and can adapt to a changing environment [2]. This collective, environmental memory interacts with the individual memories of foragers. Experienced ants, such as those in “leader-scouting” species like Lasius niger, develop internal memories of visual landmarks and routes [4, 5]. Studies show that when a conflict arises between their private memory and the social pheromone trail, experienced individuals often prioritize their own learned information, demonstrating a sophisticated interplay between individual and collective knowledge [4].

Termites, many of which are blind, also rely heavily on stigmergy for coordination [8]. They use chemical pheromones for foraging and recruitment, but also employ **sematectonic stigmergy**, where modifications to the physical structure of the nest stimulate further action [9]. For example, the deposition of a soil pellet creates a location that attracts other workers to deposit more pellets, leading to the emergent construction of complex pillars and arches. The nest structure itself becomes a persistent, physical memory of the colony's construction behavior [9]. Communication is often context-dependent; the same chemical can signal alarm at high concentrations and aggregation at low concentrations, a principle of informational efficiency [10, 11].

Honeybees utilize a more direct and symbolic form of communication. The famous **waggle dance** is a sophisticated information packet, where a returning forager communicates the distance, direction, and quality of a resource to her nestmates [12, 13]. This is not stigmergy but a form of direct, encoded messaging. The memory is first held by the individual scout, who uses path integration and landmark recognition processed in neural structures like the **mushroom bodies** [14, 15]. The dance translates this private memory into a public broadcast. Other signals, like the "tremble dance," communicate internal hive state (e.g., delays in offloading nectar), allowing the colony to regulate foraging effort based on processing capacity [13]. This represents a more centralized form of information sharing compared to the diffuse memory of an ant trail.

Social Wasps employ a communication system heavily reliant on **cuticular hydrocarbons (CHCs)**, complex chemicals on their exoskeletons [16, 17]. These CHCs form a "colony odor" that acts as a distributed memory for nestmate recognition [17]. Each wasp carries a chemical signature that is compared against an internal template learned through social contact. This allows the colony to identify and reject non-kin. Furthermore, brood (larvae and pupae) are not passive; they emit chemical signals that convey their maternity, sex, and needs, influencing worker behavior and resource allocation [18]. Some species also demonstrate social eavesdropping, learning about the fighting ability of rivals by observing them, a form of memory acquired without direct interaction [19].

3.2. Microbial Collectives (Bacterial Colonies and Slime Molds)

Even organisms without a neuron can exhibit collective memory.

Bacterial Colonies and Biofilms demonstrate environmental memory through **niche construction** [20]. As bacteria grow, they secrete an **Extracellular Polymeric Substance (EPS)**, creating a biofilm. This matrix is more than a home; it is a shared external memory and a homeostatic niche [21]. It physically traps nutrients, acts as a barrier against antibiotics, and creates chemical gradients that allow for metabolic specialization within the colony [22, 23]. The structure of the biofilm is a persistent record of the colony's history and environment. Bacteria also exhibit internal and multigenerational memory. For example, Pseudomonas aeruginosa propagates a "surface-sentient" state across generations via oscillations in intracellular cAMP levels, priming its descendants for more efficient surface adhesion [24]. This epigenetic-like inheritance ensures the collective "remembers" its sessile lifestyle.

The slime mold **Physarum polycephalum**, a single-celled amoeba with thousands of nuclei, showcases **externalized spatial memory** [25, 26]. As it forages, it leaves a trail of non-toxic extracellular slime. The organism is repelled by this slime, so the trail serves as an external memory of areas it has already explored [26, 27]. This simple mechanism allows it to solve complex mazes and optimize foraging networks by avoiding redundant paths [26]. This learned behavior—a form of habituation—can even be transferred to a "naïve" slime mold through cell fusion, demonstrating that the memory substrate is physical and distributable [28]. The morphology of its internal transport tube network also adapts to stimuli, storing localized information [28].

3.3. The Immune System

The immune system can be conceptualized as a vast, mobile multi-agent system of specialized cells that collectively defend the host. Its memory is famously associated with the adaptive immune system

(T-cells and B-cells), which creates highly specific, long-lasting memory of past pathogens. However, a more foundational form of collective memory exists in the innate immune system, known as **trained immunity** [29, 30].

Trained immunity is the functional reprogramming of innate immune cells like macrophages and Natural Killer (NK) cells [29]. After an initial encounter with a stimulus (e.g., a vaccine or infection), these cells undergo **epigenetic and metabolic rewiring** [29, 30]. Histone modifications (e.g., H3K27ac) leave key pro-inflammatory genes in a more accessible state [31]. This reprogramming does not confer specificity to a single pathogen but results in a faster, stronger, and more effective response to a subsequent, even unrelated, challenge [30]. This memory is not stored in DNA sequence but in the chromatin state, a durable yet plastic memory layer.

This memory operates at two levels. **Peripheral trained immunity** occurs in tissue-resident cells, providing localized memory at barrier sites [29, 30]. More profoundly, **central trained immunity** involves the reprogramming of **Hematopoietic Stem and Progenitor Cells (HSPCs)** in the bone marrow [29, 30]. These stem cells pass the “trained” epigenetic state to all their progeny [30]. This is a powerful mechanism for cross-generational memory transfer within the cellular lineage, ensuring that the entire population of new innate cells inherits the functional memory of past encounters, a process that can last for months or years [30].

3.4. Extended Biological Networks (Fungal Networks)

Fungal networks, particularly mycorrhizal fungi forming the “wood-wide web,” act as a biological information superhighway connecting plant root systems [32]. These extensive mycelial networks are physical conduits for transferring resources like carbon, water, and nutrients between plants. They also transmit information [32, 33]. When one plant is attacked by pests, it can send chemical distress signals through the network, inducing defensive responses in connected neighbors before they are attacked [32, 33]. The network itself, with its small-world and scale-free topology, functions as a distributed, physical memory of the forest’s connections and history [32]. Experimental studies on non-mycorrhizal fungi like Phanerochaete velutina show that mycelium can exhibit **ecological memory** [34]. Even after a physical connection to a food source is severed, the network shows a biased propensity to regrow in the direction of the former resource, suggesting that information about past growth directions is stored within the remaining network structure [34].

3.5. Migratory and Social-Learning Animal Groups

In many social animals, critical operational knowledge, such as migration routes, is not genetically innate but culturally transmitted [39, 40]. This represents a form of collective memory stored at the population level [60].

Groups of migratory birds, fish, and ungulates leverage several mechanisms to enhance navigation. The **“many wrongs” principle** suggests that by averaging the noisy navigational estimates of many individuals, the group can filter out individual errors and achieve a more accurate collective heading [35, 36]. This is a form of distributed error correction. **Leadership dynamics** also play a key role, where a small number of experienced individuals (often older) guide a larger group of naïve followers [37, 38]. This leader-follower interaction is the primary vector for **social learning**, allowing complex migration routes to be passed down through generations [37].

In ungulates like bighorn sheep and moose, migratory behavior is almost entirely cultural [39, 40]. Translocated populations introduced to new environments do not migrate initially but develop migratory patterns over decades as they explore and learn the landscape, passing this knowledge to their offspring [40]. The route itself is the “memory,” and its persistence relies on the chain of social transmission. The loss of knowledgeable older individuals can erase this cultural memory, causing the entire population to lose its migratory ability [40]. Some studies on homing pigeons suggest that this cultural “ratcheting” of efficiency can emerge from simple cognitive processes like **route averaging**, where a pair’s shared route is a weighted average of their individual preferences, rather than complex selection of the “best” route [41].

3.6. Comparative Analysis of Biological Memory Systems

The diverse strategies employed by biological systems highlight fundamental patterns in collective memory management. The following table provides a comparative summary of these systems, mapping their mechanisms and relevance to AI architecture.

System	Worker Roles / Specialization	Information Discovery & Communication	Memory Substrate & Persistence	Cross-Generation Transfer	Decay & Forgetting	Relevance to AI Architecture
Ants	Scouts, Foragers, Nurses, Soldiers [80].	Pheromone trails, Tactile cues, Tandem running [1].	**External (Environment):** Pheromone trails. Persistence depends on volatility [1].	Some evidence of route learning transfer, but primarily environmental rediscovery [4, 5].	**Pheromone evaporation:** A crucial feature for abandoning old trails [2].	Stigmergy, Ant Colony Optimization (ACO), Environment as a shared memory/blackboard [7, 9].
Honeybees	Foragers, Nurses, Guards, Cleaners (age-based polyethism) [12].	Waggle dance (symbolic), Tremble dance, Stop signal [13].	**Internal (Individual):** Neural memory (Mushroom Bodies). **Collective (Action):** The dance itself [14, 15].	Through social learning within the hive's lifespan. Queen genetics for colony traits [13].	Individual memory decay. Stop signals actively suppress outdated information [13].	Direct message passing, symbolic communication, agent-internal memory, blackboard systems [42].
Termites	Workers, Soldiers, Reproductives. Blind workers [8].	Pheromones, Vibrations, Tactile cues [8, 10].	**External (Environment):** Pheromone trails and physical nest structure (stigmergy) [9, 10].	Queen/king genetics. Trophallaxis (gut symbiont transfer) preserves digestive capability [8].	Pheromone evaporation [9].	Stigmergy, graph-based knowledge stores (the nest), context-dependent signaling [9, 11].
Bacterial Biofilms	Metabolic specialization (e.g., aerobic vs. anaerobic	Quorum sensing (chemical signals) [23].	**External (Environment):** EPS matrix (niche construction).	**Epigenetic-like inheritance:** Propagating signals	Environmental degradation, active dis-	Persistent environmental modification, shared

System	Worker Roles / Specialization	Information Discovery & Communication	Memory Substrate & Persistence	Cross-Generation Transfer	Decay & Forgetting	Relevance to AI Architecture
	layers) [22].		**Internal:** Phenotypic states [20, 21].	(e.g., cAMP oscillations) to daughter cells [24].	persal signals [22].	state management, epigenetic-like agent reprogramming [20].
Slime Molds	None (single multinucleate cell).	Chemical avoidance of its own slime [26].	**External (Environment):** Extracellular slime trail [25].	Cell fusion allows transfer of habituated state (memory) to “naïve” individuals [28].	Slime trail degradation [27].	Externalized memory, path avoidance, simple habituation models [25].
Immune System	Macrophages, NKs (Innate), T/B cells (Adaptive).	Cytokines, Chemokines, Antigen presentation.	**Internal (Cellular):** Epigenetic changes (histone modification) in trained immunity. Genomic changes in adaptive cells [29, 31].	**Central Trained Immunity:** Reprogramming of stem cells (HSPCs) passes trained state to all progeny [30].	Gradual fading of epigenetic marks over time without restimulation [30].	Agent state reprogramming, persistent functional changes, cross-generation inheritance via stem-agent “distillation” [29].
Fungal Networks	Exploratory hyphae, Transport cords.	Biochemical signals (nutrients, stress hormones) [32].	**External (Environment):** The physical mycelial network itself [32].	Spore genetics. Network persists and expands, retaining structural information [34].	Physical severing of connections, resource depletion.	Graph-based knowledge networks, physical communication infrastructure, distributed

System	Worker Roles / Specialization	Information Discovery & Communication	Memory Substrate & Persistence	Cross-Generation Transfer	Decay & Forgetting	Relevance to AI Architecture
						sensing [32].
Migratory Animals	Leaders (experienced), Followers (naïve) [37, 38].	Visual cues (following), vocalizations [37].	**Collective (Cultural):** Socially learned migration routes. Not stored in any single individual but in the group's interactions [39, 60].	**Cultural Transmission:** Social learning from parents and experienced adults to young [39, 40].	Loss of knowledgeable individuals leads to "cultural extinction" of memory [40].	Evolutionary algorithms, social learning, cross-generation memory distillation, leader-follower models [47].

4. Mapping Biological Principles to Multi-Agent AI Architectures

The biological mechanisms for collective memory provide a powerful set of analogies and design patterns for building more robust multi-agent AI systems. Translating these principles requires mapping the biological function to an appropriate computational abstraction. The core insight is that nature employs a diverse portfolio of memory strategies—from volatile environmental traces to enduring cultural knowledge—each tailored to specific ecological problems [31, 60].

4.1. Stigmergy and Environmental Memory: The World as a Whiteboard

The most direct and widely adopted biological inspiration in MAS is **stigmergy**, mirroring the behavior of ants, termites, and slime molds [9, 25]. In these systems, the environment is not a passive backdrop but an active medium for communication and memory [9].

- **Ant Colony Optimization (ACO)** algorithms are a direct implementation of ant-based stigmergy [6, 7]. Artificial ants traverse a problem space (e.g., a graph representing a logistics network) and deposit "digital pheromones" on visited paths. The intensity of this pheromone influences the path-finding decisions of subsequent agents, with evaporation mechanisms serving as the "forgetting" function [7]. This allows the swarm to collectively discover optimal solutions, such as the shortest path in a network.
- **Blackboard Systems** are an early AI architecture that conceptually mirrors stigmergy [42, 43]. A central "blackboard" acts as a shared global database where specialized agents can post hypotheses, partial solutions, and data. Other agents monitor the blackboard and contribute based on their expertise, reacting to the changes made by others [42]. This is analogous to termites building a nest; the structure on the blackboard (or in the nest) stimulates a new round of action. Modern implementations use this for complex coordination, where agents asynchronously contribute to a shared problem-solving space [43, 61].

- **Market-Based Allocation** can be seen as a form of stigmergy where the “traces” are economic signals [44]. Agents “bid” on tasks, and the prices in the virtual market serve as an environmental signal of demand and resource availability. Other agents react to these price signals to make local decisions, leading to emergent, system-wide resource allocation [44, 45].

4.2. Direct Communication and Individual Cognition: The Honeybee Model

While stigmergy excels at diffuse coordination, it is inefficient for communicating precise, complex information. Honeybees solve this with the waggle dance, a form of direct, symbolic communication [13].

- **Message Passing Architectures** in MAS are the computational equivalent. Agents explicitly send structured messages to other agents or to a broadcast channel [62]. This allows for the transmission of rich, specific information—such as a vector, a coordinate, or a block of code—far more efficiently than modifying a shared environment. This is suitable for tasks where agents need to share high-fidelity private knowledge.
- **Shared Memory Architectures** in modern MAS, especially those using **Retrieval-Augmented Generation (RAG)**, reflect the interplay between an agent’s internal knowledge (analogous to the bee’s neural memory) and shared information [59]. An agent queries a shared vector database (a form of collective memory) to augment its internal context (its prompt) before generating a response. This allows individual agents to leverage a vast, persistent knowledge base that they do not possess internally [59, 63].

4.3. Epigenetics and Cultural Transmission: Memory Across Generations

Biological systems ensure knowledge persists beyond the lifespan of an individual. This is achieved through mechanisms like epigenetic inheritance and cultural transmission, which have powerful analogues in AI [29, 40].

- **Evolutionary Algorithms (EAs)** directly mimic the process of natural selection and inheritance [47]. A population of agents (or agent configurations) is evaluated against a fitness function. The most successful agents are “selected” to “reproduce,” creating a new generation that inherits their traits, often with mutation and crossover [47]. This is a direct parallel to the cultural evolution of migration routes, where successful strategies are passed down and refined over generations [40].
- **Cross-Generation Memory Transfer** can be explicitly designed. Drawing inspiration from the immune system’s central trained immunity, a “parent” generation of agents that has completed a task can have its collective memory **distilled** [30]. An orchestrator or validator agent can identify the most critical, high-utility memories (e.g., proven facts, successful tool-use sequences, key environmental features) from the archival record of the old generation. This curated knowledge base can then be used to initialize the memory of a new “child” generation, giving them a significant head start and preventing them from having to re-learn everything from scratch [47]. This is analogous to how HSPCs endow all future immune cells with a “trained” phenotype [30].
- **Multi-Agent Reinforcement Learning (MARL)** frameworks can incorporate lineage. As agents learn and evolve strategies, the successful policies can be stored and passed on [46]. For example, research into the Iterated Prisoner’s Dilemma using MARL shows how dominant, cooperative strategies can emerge and persist in a population over time, which is a form of learned, cultural memory within the agent society [64].

The following table maps these biological inspirations to their corresponding AI architectural concepts.

Biological Principle	Biological Example(s)	AI Architectural Analogue	Key Functionality
Stigmergy / Externalized Memory	Ant pheromone trails, Termite nests, Slime mold slime [1, 9, 25].	Ant Colony Optimization (ACO), Particle Swarm Optimization (PSO), Blackboard Systems, Digital Pheromones [7, 42, 65].	Indirect coordination, environmental state as shared memory, emergent pathfinding and optimization.
Niche Construction	Bacterial biofilms creating an EPS matrix [20, 21].	Persistent Environmental Modification, Shared Causal Memory.	Agents alter the shared environment, creating state that influences future agents' behavior and decisions.
Direct / Symbolic Communication	Honeybee waggle dance [13].	Agent-to-Agent Message Passing, Shared Episodic/Semantic Memory [62].	High-fidelity transmission of complex, specific information (vectors, commands, observations).
Epigenetic Reprogramming	"Trained Immunity" in innate immune cells [29, 30].	Agent State Reprogramming, Shared Policy Memory.	Persistently altering an agent's functional parameters (e.g., heuristics, retrieval weights) based on experience, creating a long-term "trained" state.
Centralized Inheritance	Reprogramming of Hematopoietic Stem Cells (HSPCs) [30].	Cross-Generation Memory Distillation, Evolutionary Algorithms [47].	A parent generation's knowledge is curated and inherited by a new generation of "stem" agents, ensuring durable knowledge transfer.
Social Learning / Cultural Transmission	Migration routes in ungulates and birds (Leader-Follower) [37, 40].	Evolutionary Algorithms, Social Learning in MARL, Lineage Tracking [46, 47].	Knowledge is transferred from experienced agents to naïve agents through observation or direct guidance, allowing

Biological Principle	Biological Example(s)	AI Architectural Analogue	Key Functionality
			for knowledge accumulation over time.
Distributed Sensing / Error Correction	"Many Wrongs" principle in flocks and shoals [35, 36].	Ensemble Methods, Voting Mechanisms in MAS.	Aggregating noisy inputs from multiple agents to produce a more accurate collective output; error reduction through consensus.

5. A Governed Multi-Agent Memory Architecture

Drawing from the rich tapestry of biological memory, we propose a concrete architecture for a governed multi-agent memory system. The central design philosophy is the separation of **free internal cognition** from **externally governed action** [53]. Agents are free to reason, hypothesize, and maintain private beliefs. However, any action that mutates the shared state—writing to collective memory, calling an external tool, or communicating with another agent—is a governed transaction that must be authorized, verified, and logged with immutable provenance [50, 53]. This architecture is designed to make the collective memory auditable, resilient, and trustworthy [48].

5.1. Architectural Overview and Core Principles

The proposed architecture is built upon three foundational pillars:
1. **Structured, Verifiable Memory:** All memories are not created equal. Information is categorized and stored in structured records enriched with metadata for provenance, lineage, and relevance. All entries are cryptographically signed to ensure integrity [55].
2. **Policy-Driven Governance:** A dedicated governance layer intercepts all attempts to modify shared state [53]. It enforces rules based on agent identity, data provenance, and task context, acting as the "immune system" of the collective [56].
3. **Inspectable Lifecycle:** The entire lifecycle of a memory—from its creation to its archival or deletion—is tracked [55]. This creates a "chain of evidence" that allows human auditors and automated systems to reconstruct the reasoning behind any agent action [49, 50].

This approach translates directly from biological principles. The "receipt" of a governed action is akin to a stable pheromone trail—a public, verifiable record of a past success. An agent's internal thought process is like a honeybee's private path integration—ephemeral and internal—but the waggle dance is a public, structured broadcast subject to the colony's interpretation and feedback [14].

5.2. Memory Structure and Types

To move beyond undifferentiated vector stores, the system utilizes a multi-layered memory structure with distinct types, each serving a specific cognitive function.

- **Episodic Memory:** Stores a raw, chronological log of events and experiences. This includes sensor readings, raw tool outputs, and direct inter-agent messages. It is the agent's "lived experience" and serves as the ultimate ground truth for what has occurred. These records are write-once and immutable.
- **Semantic Memory:** Stores distilled facts, concepts, and knowledge extracted from episodic memory or provided by external sources. This is the agent's "knowledge base" of general prin-

ciples about the world (e.g., "API key X is used for service Y"). These records are versioned to allow for updates as knowledge evolves.

- **Procedural Memory:** Stores successful sequences of actions or "skills." For example, a proven workflow for deploying a web service. This memory is crucial for specialization, allowing agents to learn and reuse effective behaviors. This is analogous to a forager ant memorizing a successful route [4].
- **Causal Memory:** A more advanced memory type that explicitly links cause-and-effect relationships. It stores records in the form of "If action A is taken in state S, outcome O is likely." This memory is built from analyzing correlations in the episodic log and is critical for planning and counterfactual reasoning.
- **Archival Memory:** A long-term, compressed storage for memories that are no longer immediately relevant but may have historical value [52]. This prevents the active memory from being cluttered while preserving a complete record for audit or future, deep analysis.

5.3. The Memory Record Schema: Metadata is Key

The power of this architecture lies not in the content of the memory but in its metadata. Every memory record, regardless of type, conforms to a standard schema that enables governance and advanced retrieval [48, 55].

```
{
  "RecordID": "unique_content_hash_blake3", // Content-addressable ID
  "ParentID": "id_of_source_record",        // For distilled/summarized memories
  "Timestamp": "2026-06-21T14:30:00Z",       // Creation time
  "AuthorAgentID": "agent_xyz_123",          // ID of the writing agent
  "MemoryType": "Semantic",                  // Episodic, Semantic, Procedural, Causal
  "Content": { ... },                       // The actual memory data
  "ProvenanceID":
"receipt_or_source_id",    // Link to the receipt or user input that created this
memory
  "LineageGraph": ["id1", "id2", ...],       // DAG of parent memories used in its
creation
  "RelevanceScore": {                       // Multi-dimensional score
    "base_importance": 0.8,
    "last_access": "2026-06-21T15:00:00Z",
    "access_count": 5
  },
  "DecayState": {                           // For spaced repetition scheduling
    "stability": 15.0, // In days
    "difficulty": 0.3
  },
  "ArchivalStatus": "Active",                // Active, Buffered, Archived
  "VerificationSignature": "ed25519_signature" // Signature of the record by the
author agent
}
```

This schema turns the memory from a simple database into a governable, auditable ledger [55]. The `ProvenanceID` and `LineageGraph` are particularly critical [48]. They create a directed acyclic graph (DAG) tracing every piece of derived knowledge back to its source—either a direct user input or a verifiable receipt from an external action [48, 56]. This is the foundation for trusting retrieved information.

5.4. The Memory Lifecycle: From Ingestion to Archival

The lifecycle of memory within this architecture is an active, managed process [55].

1. **Write (Ingestion):** An agent proposes a write to memory. This proposal is intercepted by the governance layer. The proposal must be justified by citing its source (e.g., a direct observation, a tool output, or a set of parent memories) [53].
2. **Store (Indexing):** If the write is approved, a new memory record is created. It is signed by the agent, and its `RecordID` is generated via a content hash. The record is then indexed for retrieval (e.g., via vector embedding of its content and keyword extraction from its metadata).
3. **Retrieve:** Agents query the memory system with a task-specific need. The retrieval engine uses a multi-signal scoring model (detailed below) to return the most relevant and trustworthy records [52].
4. **Execute:** When an agent proposes a consequential external action (e.g., an API call), it must submit a **"justification proof"** to the governance layer [50, 53]. This proof contains the agent's intent and the `RecordID`s of the memories supporting its decision. The governance layer verifies this proof against its policies. If approved, the action is executed.
5. **Share (Receipt Generation):** Upon successful execution, a cryptographic **"receipt"** is generated [50, 66]. This receipt is an immutable `Episodic` memory record detailing the action taken, its parameters, and its outcome. It is linked to the `ProvenanceID` of the memories that justified it, closing the loop of accountability [49].
6. **Forget & Rollback:** Memory records are subject to a decay process. If a record's relevance score drops below a certain threshold and it has not been accessed for a long period, its status is changed to `Archived` [52]. An explicit "forget" command can be issued (e.g., for privacy compliance), which either cryptographically shreds the record or marks it as non-retrievable, while preserving its ID in lineage graphs for audit integrity [56].

5.5. Retrieval Strategies: Multi-Signal Scoring

Simple semantic search is insufficient. The retrieval engine must act as a gatekeeper, prioritizing useful and trustworthy information. To achieve this, retrieval is based on a composite scoring model that combines multiple signals:

- **Semantic Relevance:** The cosine similarity between the query embedding and the memory content embedding. This is the baseline for "what" the memory is about.
- **Temporal Relevance:** A decay function that down-weights older memories, but with a boost for recently accessed items [52]. This prioritizes recency.
- **Utility Score:** A measure of how often a memory has been part of a successful justification proof [67]. Memories that consistently lead to successful actions are ranked higher. This is a form of reinforcement.
- **Provenance Trust:** A score based on the origin of the memory [55]. Memories created directly from user input or from verifiable receipts have the highest trust. Memories derived from other memories, especially those created by less reliable agents or with long lineage chains, have lower trust.
- **Task Context:** The system can boost the scores of certain `MemoryType`s based on the task at hand. For a planning task, `Causal` and `Procedural` memories might be prioritized.

This multi-signal approach ensures that the memories surfaced to an agent are not just topically related but are also timely, trustworthy, and historically useful.

6. Operational Models for the Governed Memory System

The architectural framework outlined above requires concrete operational models to function effectively. This section details the proposed models for scoring and relevance decay, cross-generational memory transfer, and the practical implementation of provenance-based governance.

6.1. Relevance, Decay, and Archival Scoring Model

To prevent stale or low-utility information from overwhelming an agent's context window, we propose a dynamic scoring model that governs a memory record's lifecycle. A record's final **Retrieval Score** for a given query is a weighted sum of several factors, ensuring a holistic assessment of its value.

The **Retrieval Score** **`S(m, q, t)`** for a memory record `m`, given a query `q` at time `t`, is calculated as:

`S(m, q, t) = w_sim * Sim(m, q) + w_temp * Temp(m, t) + w_util * Util(m) + w_prov * Prov(m)`

Where:
* `w_sim`, `w_temp`, `w_util`, `w_prov` are configurable weights that sum to 1.
* **`Sim(m, q)`** **(Semantic Similarity):** The standard cosine similarity between the vector embeddings of the memory's content and the query. This score ranges from -1 to 1.
* **`Temp(m, t)`** **(Temporal Score):** This score balances recency and frequency, inspired by spaced repetition algorithms [51]. It is calculated based on the memory's `DecayState` metadata. We can adapt a model like the Free Spaced Repetition Scheduler (FSRS), where a memory's "retrievability" (R) at time `t` is a function of its "stability" (S), which represents how long the memory is likely to be retained [52]. `R = e^(-Δt / S)`, where `Δt` is the time since last access. The `Temp(m, t)` score is this retrievability value, ranging from 0 to 1. When a memory is accessed, its stability increases, fighting decay. This is more nuanced than simple linear decay.
* **`Util(m)`** **(Utility Score):** This score is a measure of a memory's historical usefulness. It is a normalized value derived from the `access_count` metadata, potentially with a logarithmic scaling to provide diminishing returns: `Util(m) = log(1 + m.metadata.access_count)`. More importantly, this count should be incremented only when the memory is part of a justification proof for a successful action, creating a reinforcement loop [67].
* **`Prov(m)`** **(Provenance Score):** This score reflects the trustworthiness of the memory's origin [55]. It is a value between 0 and 1 assigned based on a trust hierarchy:
* `1.0`: Direct user input or a cryptographically verified receipt of an external action. This is ground truth [50, 53].
* `0.8`: A semantic memory distilled from a high-trust episodic record.
* `0.6`: A procedural memory derived from a sequence of successful actions.
* `0.4`: A causal memory inferred by an agent, as this is speculative.
* `0.2`: A memory received from another agent with a low trust score.

Archival and Forgetting: The archival process is driven by this scoring system. Inactive memories do not generate a retrieval score. A background process periodically evaluates the `DecayState` of all active memories. If a memory's retrievability `R` drops below a low threshold (e.g., 0.1) and it has not been accessed for a significant period, its `ArchivalStatus` is changed from 'Active' to 'Archived' [52]. Archived memories are moved to cheaper, slower storage and are excluded from standard retrieval queries, but they remain available for deep audits via their `RecordID`.

6.2. Cross-Generation Transfer: Inheritance and Distillation

Inspired by cultural transmission in animal groups and central trained immunity, this model provides a mechanism for knowledge to persist beyond the lifecycle of a single agent or task force [30, 39, 40]. This is crucial for creating learning organizations of agents that accumulate wisdom over time.

The process involves three stages:
1. **Selection:** At the conclusion of a major project or after a fixed "generation" period, an automated "Curator" agent is initiated. This agent's task is to review the `Archival Memory` of the completed generation. It identifies the most valuable memories by querying for records with the highest final `UtilityScore` and `ProvenanceScore` [67, 55]. This selects for memories that were both demonstrably useful and highly trustworthy.
2. **Distillation:** The Curator agent then performs a distillation process [47]. It may select a variety of memory types:
* **High-value `Semantic` records** (critical facts) are preserved as is.
* **High-value `Procedural` records** (successful workflows) are compiled into a "best practices" library.
* Clusters of related `Episodic` records that led to a key outcome are summarized into a new, high-level `Causal` or `Semantic` memory.
3. **Inheritance:** This curated and distilled set of memories forms the **"Foundation Knowledge Base"** for the next generation of agents [47]. When new agents are instantiated, their memory is pre-populated with this foundational set. This is analogous to young ungulates being born into a herd that already knows the migration route; they don't start from a blank slate [39, 40]. The `LineageGraph` of these inherited memories would trace back to the Curator agent and the previous generation, maintaining a continuous thread of provenance [48].

6.3. Provenance and Governance in Action

This model operationalizes the "free cognition, governed action" principle by embedding verification into the agentic loop [53].

Let's walk through an example: An agent is tasked with ordering more computing resources for a project.
1. **Cognition (Internal):** The agent queries its memory. It retrieves a `Semantic` memory stating "Project 'Blue' compute budget is $500/day," an `Episodic` memory showing current usage is at $480/day, and a `Procedural` memory detailing the API call to provision new servers.
2. **Intent (Proposal):** The agent formulates an intent: "Provision one new 'm5.large' server for Project 'Blue'." It packages this intent with the `RecordID`s of the three memories it used for its reasoning. This package is the **"justification proof."** [50, 53].
3. **Governance (Verification):** The proof is sent to the governance layer (the Sovereign Execution Broker) [53]. The broker is a separate, high-privilege service. It executes a verification policy:
* It checks the `VerificationSignature` on all cited memory records to ensure they haven't been tampered with [55].
* It traces the `ProvenanceID` of the budget memory back to the original user command that set the budget, confirming its authority [48].
* It executes a real-time check against an external billing API (a "live-state drift" check) to confirm the current usage is indeed $480/day [53].
* It confirms the proposed action (provisioning a server) is within the policy for Project 'Blue'.
4. **Action (Execution):** The verification passes. The governance layer generates temporary, single-use credentials and executes the API call to provision the server [53]. The agent itself never holds the long-term credentials.
5. **Record (Receipt):** The API call returns a success message with the new server's ID. The governance layer creates a new, immutable `Episodic` memory record—the **receipt** [50]. This receipt contains the action taken, its parameters, the successful outcome, and the hash of the justification proof. It is signed by the governance system itself, making it a non-repudiable record of a verified action [49, 66].

This closed-loop system ensures that every consequential action is justified, verified, and recorded. The memory system becomes an auditable ledger, not just a passive store of data [50].

7. Translating Principles into Governed System Constraints

The proposed architecture is not merely a technical specification but a manifestation of a set of core principles derived from observing stable biological collectives. To ensure a multi-agent system operates in a robust and governable manner, these principles must be enforced as explicit system constraints.

- **Free Internal Cognition, Externally Governed Actions:** This is the foundational constraint. Agents must be allowed to perform internal reasoning, planning, and simulation without constant oversight. This “cognitive freedom” is essential for creativity and complex problem-solving. However, the moment an agent attempts to affect the shared world—by writing to collective memory, communicating with another agent, or using a tool—its action must pass through a governance checkpoint [53, 54]. This mirrors a wasp colony where an individual can forage freely, but its acceptance back into the nest depends on passing a chemical “odor check” [17].

- **Consequential Actions Must Carry Evidence and Provenance:** No action that consumes resources or has external effects should be permitted without a “justification proof” [50]. This proof is a structured artifact linking the agent’s intent to the specific memory records that informed its decision. Upon execution, the system must generate a non-repudiable “receipt” [50, 66]. This constraint eliminates “because I said so” reasoning from agents and ensures every significant event is causally anchored to a verifiable chain of evidence [49].

- **Memories Do Not Outrank Receipts or Causal Records:** An agent’s memory represents its beliefs, which can be flawed or outdated. A receipt, however, is a verifiable record of a past event. The retrieval and governance systems must be configured to prioritize receipts and other high-provenance records over an agent’s own derived or semantic memories, especially in cases of conflict [50, 53]. An agent cannot be allowed to “believe” a budget is $1000 if a verifiable receipt shows it was set to $500. This maintains the integrity of ground truth within the collective.

- **Lineage Must Be Preserved Across Generations:** When knowledge is distilled and inherited by new agents, its provenance must not be erased [48]. The new, foundational memories given to a “child” agent must contain lineage metadata that traces back to the “parent” generation’s records from which they were derived [56]. This prevents knowledge laundering and ensures that even inherited beliefs can be audited back to their original empirical source, mimicking the unbroken chain of cultural transmission [40].

- **Worker Specialization Improves Cognition Without Ungoverned Execution:** Agents should be encouraged to specialize, developing deep procedural and semantic memory in specific domains, much like the role differentiation in an ant colony [68, 80]. However, specialization must not grant an agent the ability to bypass governance. A specialized “database agent” may be highly efficient at writing SQL queries, but its ability to execute those queries against a production database remains subject to the same external verification as any other agent [57]. Specialization enhances the quality of proposals, not the authority to execute them.

By building these constraints into the very fabric of the multi-agent system, we can foster a collective that balances the autonomy needed for intelligent adaptation with the structural governance required for safety and accountability [53, 54].

8. Conclusion

The formidable challenge of building scalable, trustworthy multi-agent AI systems compels us to look beyond conventional computing paradigms and seek inspiration from nature’s most successful decentralized collectives [76]. Biological swarms, from ant colonies to the human immune system, have spent millennia perfecting the art of collective memory, balancing individual autonomy with collective stability through intricate systems of communication, governance, and knowledge transfer [31]. This report has demonstrated that these biological principles are not mere philosophical curiosities; they are actionable design patterns for engineering the future of artificial intelligence.

Our survey revealed a rich portfolio of natural memory mechanisms: the dynamic, self-forgetting environmental memory of stigmergic insects [1, 9]; the precise, symbolic communication of honeybees [13]; the robust, non-specific epigenetic memory of trained immunity [29]; and the durable, culturally transmitted knowledge of migratory herds [39]. By mapping these strategies to AI constructs, we have highlighted a path forward that moves beyond the limitations of today's brittle and opaque memory systems [55, 59].

The primary contribution of this report is the proposal for a **Governed Multi-Agent Memory Architecture**, a concrete framework designed to bring accountability and resilience to artificial collectives. Its core tenets—structured memory with cryptographic provenance, multi-signal retrieval scoring that values trust as much as relevance, a governed lifecycle where actions require justification, and a model for cross-generational knowledge distillation—directly address the critical failure modes of current systems [55, 56]. By enforcing a clear separation between an agent's unregulated internal cognition and its verifiably governed external actions, this architecture creates a system that is both intelligent and inspectable [53].

The implementation of such a system is the clear next step. Future research should focus on building and testing prototypes of this architecture, refining the proposed scoring and decay models, and exploring the emergent behaviors of agent collectives operating under these governance constraints. The ultimate goal is to create multi-agent systems that can learn, adapt, and collaborate effectively over long time horizons, while providing the non-repudiable audit trails necessary for deployment in high-stakes domains [49, 50]. By learning from the hive, we can build more intelligent and more trustworthy hard drives, ushering in an era of accountable artificial swarms.

9. References

1. ANT COMMUNICATION AND COGNITION - academia.edu (https://www.academia.edu/3133950/ANT_COMMUNICATION_AND_COGNITION)
2. Temperature limits trail following behaviour through pheromone decay in ants - researchgate.net (https://www.researchgate.net/publication/51752893_Temperature_limits_trail_following_behaviour_through_pheromone_decay_in_ants)
3. Ant foraging on complex trails: route learning and the role of trail pheromones in Lasius niger - journals.plos.org (https://journals.plos.org/plosone/article?id=10.1371%2Fjournal.pone.0149720)
4. Memory and chemical communication: the orientation of two mass-recruiting ant species - researchgate.net (https://www.researchgate.net/publication/225805008_Memory_and_chemical_communication_the_orientation_of_two_mass-recruiting_ant_species)
5. Ant foraging on complex trails: route learning and the role of trail pheromones - journals.biologists.com (https://journals.biologists.com/jeb/article/216/2/188/11672/Ant-foraging-on-complex-trails-route-learning-and)
6. Collective Stigmergic Optimization - arxiv.org (https://arxiv.org/pdf/cs/0403001)
7. Ant-algorithms and stigmergy - unibo.it (https://lia.disi.unibo.it/courses/2006-2007/PSI-LS/pdf/roli/dorigo2000-ant_algorithms_and_stigmergy.pdf)
8. Biology / Behavior of Eastern Subterranean Termites - ncsu.edu (https://content.ces.ncsu.edu/biology-behavior-of-eastern-subterranean-termites)
9. Stigmergy - en.wikipedia.org (https://en.wikipedia.org/wiki/Stigmergy)

10. A Review of Termite Pheromones: Multifaceted, Context-Dependent and Rational Chemical Communications - frontiersin.org (https://www.frontiersin.org/journals/ecology-and-evolution/articles/10.3389/fevo.2020.595614/full)
11. A Review of Termite Pheromones: Multifaceted, Context-Dependent, and Rational Chemical Communications - researchgate.net (https://www.researchgate.net/publication/348381113_A_Review_of_Termite_Pheromones_Multifaceted_Context-Dependent_and_Rational_Chemical_Communications)
12. The Colony and Its Organization - canr.udel.edu (https://canr.udel.edu/maarec/honey-bee-biology/the-colony-and-its-organization/)
13. On the Communication of Honeybee Swarms - springer.com (https://link.springer.com/article/10.1007/s00265-022-03218-1)
14. Learning and memory in honeybees: from behavior to neural substrates - fu-berlin.de (https://www.bcp.fu-berlin.de/biologie/arbeitsgruppen/neurobiologie/ag_menzel/publications/Res/Me-Review-nature-Reviews-Neurosc-2012.pdf)
15. Learning and Memory in Honeybees - learnmem.cshlp.org (https://learnmem.cshlp.org/content/8/2/53.full)
16. Social Wasps - genent.cals.ncsu.edu (https://genent.cals.ncsu.edu/bug-bytes/social-insects/social-wasps/)
17. Plasticity of Nestmate Recognition Cues in Social Wasps: A Review and Future Perspectives - frontiersin.org (https://www.frontiersin.org/journals/ecology-and-evolution/articles/10.3389/fevo.2019.00444/full)
18. The role of brood in social insect societies - sciencedirect.com (https://www.sciencedirect.com/science/article/pii/S2666515825000022)
19. Social wasps “eavesdrop” to learn about the fighting ability of rivals - cell.com (https://www.cell.com/current-biology/pdf/S0960-9822(20)30740-5.pdf)
20. Building the microbiome in health and disease: Niche construction, adaptation, and maladaptation - pmc.ncbi.nlm.nih.gov (https://pmc.ncbi.nlm.nih.gov/articles/PMC4528496/)
21. Biofilms as a microbial homeostatic niche that synchronizes bacterial communities - nature.com (https://www.nature.com/articles/s41579-022-00791-0)
22. Bacterial biofilms: Home of a persistent infection - pmc.ncbi.nlm.nih.gov (https://pmc.ncbi.nlm.nih.gov/articles/PMC10305407/)
23. Bacterial biofilms: Home of a persistent infection - rsc.org (https://pubs.rsc.org/en/content/articlehtml/2025/pm/d5pm00094g)
24. Multigenerational memory and extinction of a bacterial collective behavior - pnas.org (https://www.pnas.org/doi/10.1073/pnas.1720071115)
25. Externalized memory in slime mould - philsci-archive.pitt.edu (https://philsci-archive.pitt.edu/20082/7/Externalized%20memory%20in%20slime%20mould%20-%20REVISION%20.pdf)
26. Slime mold uses externalized spatial memory to navigate in complex environments - pnas.org (https://www.pnas.org/doi/full/10.1073/pnas.1215037109)
27. Slime mold uses externalized spatial memory to navigate in complex environments - pubmed.ncbi.nlm.nih.gov (https://pubmed.ncbi.nlm.nih.gov/23045640/)
28. site:pmc.ncbi.nlm.nih.gov Physarum polycephalum memory habituation slime mold review - pmc.ncbi.nlm.nih.gov (https://pmc.ncbi.nlm.nih.gov/articles/PMC10667701/)
29. Trained immunity: A memory for all seasons - sciencedirect.com (https://www.sciencedirect.com/science/article/abs/pii/S0952791525001062)
30. Trained Immunity: A Memory for All Seasons - nature.com (https://www.nature.com/articles/s41577-020-0285-6)

31. Collective intelligence from bacteria to humans - nature.com (https://www.nature.com/articles/s41467-023-43099-0)
32. Mycorrhizal Networks Facilitate Tree Communication, Learning, and Memory - link.springer.com (https://link.springer.com/chapter/10.1007/978-3-319-75596-0_10)
33. The “wood wide web” of forests is a real-life phenomenon - nph.onlinelibrary.wiley.com (https://nph.onlinelibrary.wiley.com/doi/full/10.1002/ppp3.10178)
34. Ecological memory of mycelial networks. - nature.com (https://www.nature.com/articles/s41396-019-0536-3)
35. Group navigation and the ‘many-wrongs principle’ in models of animal movement - academia.edu (https://www.academia.edu/722355/Group_navigation_and_the_many_wrongs_principle_in_models_of_animal_movement)
36. Group navigation and the many-wrongs principle: an emergent property of migrating animals - springer.com (https://link.springer.com/article/10.1186/s40462-016-0083-8)
37. Collective animal navigation and migratory culture: from theoretical models to empirical evidence - researchgate.net (https://www.researchgate.net/publication/324016203_Collective_animal_navigation_and_migratory_culture_from_theoretical_models_to_empirical_evidence)
38. Leadership in fish shoals - researchgate.net (https://www.researchgate.net/publication/263154707_Leadership_in_fish_shoals)
39. The cultural basis of ungulate migration - science.org (https://www.science.org/doi/10.1126/science.aat0985)
40. Is ungulate migration culturally transmitted? Evidence of social learning from translocated animals - researchgate.net (https://www.researchgate.net/publication/327479687_Is_ungulate_migration_culturally_transmitted_Evidence_of_social_learning_from_translocated_animals
41. Beyond Average: The Role of Forgetting and Perspective-Taking in the Cultural Evolution of Homing Pigeon Routes - elifesciences.org (https://elifesciences.org/reviewed-preprints/108054)
42. A Survey of Multi-Agent Systems and Workflow: Design Principles and Implementation Strategies for Modern AI Systems - researchgate.net (https://www.researchgate.net/publication/395128299_Multi-Agent_Systems_and_Workflow_Survey_Design_Principles_and_Implementation_Strategies_for_Modern_AI_Systems
43. LbMAS: a Large Language Model based Multi-Agent System with the Blackboard Architecture for Complex Task Solving - arxiv.org (https://arxiv.org/html/2507.01701v1)
44. Market-inspired Mission Economies for Agent-based AIs - arxiv.org (https://arxiv.org/html/2502.02311v3)
45. Market-Based Task Allocation Mechanisms in Multiagent Systems under Soft and Hard Real-Time Constraints - acm.org (https://dl.acm.org/doi/fullHtml/10.1145/3356467)
46. Multi-Agent Reinforcement Learning: A Methodological Survey, Analysis and Taxonomy - arxiv.org (https://arxiv.org/html/2407.17030v1)
47. EvoMAS: Evolutionary Generation of Multi-agent Systems - arxiv.org (https://arxiv.org/html/2602.06511v1)
48. MemLineage: A Verifiable, Lineage-Guided Defense against Memory Poisoning in Autonomous Agents - arxiv.org (https://arxiv.org/html/2605.14421)
49. Computable Responsibility: A Structural Approach to Safe & Auditable Reinforcement Learning for Multi-Agent AI - arxiv.org (https://arxiv.org/html/2605.17169)
50. Agent Execution Record (AER): A Structured Representation of Agent Trajectories for Comprehensive Evaluation - arxiv.org (https://arxiv.org/html/2603.21692)
51. Learning to Remember: A Biologically-Inspired Method for Continual Learning with Spaced Rehearsal - arxiv.org (https://arxiv.org/abs/2506.12034)

52. AgentMemory: A Spaced-Repetition-Based Method for Enhancing the Long-Term Memory of AI Agents - arxiv.org (https://arxiv.org/html/2603.18330)
53. Sovereign Execution Brokers: A Structural Approach to Governance in Agentic AI - arxiv.org (https://arxiv.org/html/2606.20520v1)
54. Distributed Trust Frameworks - A Structural Approach to AI Safety in High-Stakes Automation - arxiv.org (https://arxiv.org/html/2605.15228)
55. MemArchitect: A Fine-grained, Field-Level, Governance-Oriented Format for Autonomous Agent Memory - arxiv.org (https://arxiv.org/html/2606.04990)
56. Mnemonic Sovereignty: Towards Verifiable Memory Governance in Autonomous Agents - arxiv.org (https://arxiv.org/html/2604.16548v2)
57. AgentNet: A Scalable and Verifiable Multi-Agent System with Distributed Coordination - arxiv.org (https://arxiv.org/html/2504.00587.pdf)
58. Emergent Collective Memory in Decentralized Multi-Agent AI Systems - researchgate.net (https://www.researchgate.net/publication/398602345_Emergent_Collective_Memory_in_Decentralized_Multi-Agent_AI_Systems)
59. A Survey and Roadmap for LLM-based Autonomous Agents - arxiv.org (https://arxiv.org/html/2601.12560v1)
60. Collective Memory in Animal Groups - wires.onlinelibrary.wiley.com (https://wires.onlinelibrary.wiley.com/doi/10.1002/wcs.70035)
61. Information discovery in data lakes with a large language model based multi-agent system - arxiv.org (https://arxiv.org/html/2510.01285v1)
62. A Survey on Communication in Multi-Agent Deep Reinforcement Learning - springer.com (https://link.springer.com/article/10.1007/s10458-023-09633-6)
63. Guide a guide to building multi-agent systems with shared memory - vectorize.io (https://hindsight.vectorize.io/guides/2026/04/21/guide-building-multi-agent-systems-with-shared-memory)
64. Learning a Dominant Bilateral Reciprocity Strategy for the Repeated Game of Prisoner's Dilemma - nature.com (https://www.nature.com/articles/s41467-025-67178-6)
65. Swarm intelligence: a new C-paradigm for social and behavioral sciences? - arxiv.org (https://arxiv.org/pdf/nlin/0502057)
66. PROV-AGENT: a provenance model for workflows with AI agent interactions - arxiv.org (https://arxiv.org/html/2508.02866v1)
67. Cognitive Weighting: An Online Learning Framework for Context Generation - arxiv.org (https://arxiv.org/html/2605.03675v1)
68. AOrchestra: An Adaptive Orchestrator for Multi-Agent Systems - arxiv.org (https://arxiv.org/html/2602.03786v2)
69. A new type of aggregation pheromone in the Virginia termite, Reticulitermes virginicus (Banks) (Blattodea: Rhinotermitidae) - nature.com (https://www.nature.com/articles/s41598-024-73310-1)
70. Termite social communication: an under-researched area of insect social biology - wiley.com (https://esj-journals.onlinelibrary.wiley.com/doi/10.1111/1440-1703.12510)
71. Social Organisation and Information Transfer in Schooling Fish - researchgate.net (https://researchgate.net/publication/229868541_Social_Organisation_and_Information_Transfer_in_Schooling_Fish)
72. site:pmc.ncbi.nlm.nih.gov trained immunity epigenetic memory review - pmc.ncbi.nlm.nih.gov (https://pmc.ncbi.nlm.nih.gov/articles/PMC6121175/)
73. site:pmc.ncbi.nlm.nih.gov Physarum polycephalum memory habituation slime mold review - pmc.ncbi.nlm.nih.gov (https://pmc.ncbi.nlm.nih.gov/articles/PMC6553583/)

74. A Survey on Communication in Multi-Agent Deep Reinforcement Learning - arxiv.org (https://arxiv.org/abs/2203.08975)
75. Self-Revising Discovery Systems: A Categorical Framework for Autonomy in Science - arxiv.org (https://arxiv.org/pdf/2604.15034)
76. SwarmSys: A Swarm Intelligence-inspired Paradigm for Multi-agent Coordination - arxiv.org (https://arxiv.org/html/2510.10047v1)
77. Agentic AI Systems: A Survey and Roadmap - arxiv.org (https://arxiv.org/html/2509.10147)
78. Multi-Agent Systems: A Survey from a Machine Learning Perspective - cmu.edu (https://www.cs.cmu.edu/~mmv/papers/MASsurvey.pdf)
79. On the study of insect societies - ebsco.com (https://www.ebsco.com/research-starters/science/insect-societies)
80. Ant’s Personality and Its Dependence on Foraging Styles - antwiki.org (https://www.antwiki.org/w/images/a/a6/Reznikova,_Z._(2021)_Ants%E2%80%99_Personality_and_Its_Dependence_on_Foraging_Styles_(10.3389@fevo.2021.661066).pdf)